



# Great Expectations & The Wonderful World of Data Quality Tools in Python

Sam Bail @spbail, Data Umbrella, March 2021


# Hi, I'm Sam! (@spbail)

- I'm originally from Germany, currently based in NYC
- I have a PhD in semantic web technologies with a focus on data representation formalisms (Linked Data, OWL, RDF, in case that rings a bell...)
- I've been doing "data work" for a while, mostly working with 3rd party healthcare data
- And now I'm an engineering director at Superconductive, the core maintainers behind Great Expectations

# Agenda

- Part 1: The Wonderful World of (Open Source) Data Quality in Python
  - Types of “data quality” tools
  - Overview of some prominent ones
- Part 2: Great Expectations: Overview and motivation
- Part 3: “Getting started” live demo of Great Expectations
- Q&A

# The challenge: Data workflows today are a mess



- Data pipelines are **brittle** and often **fail**, both loudly and silently
- **Tacit knowledge** scattered among domain experts, technical experts, and the code and data itself
- **Maintenance** is time-consuming, expensive, and morale-killing
- **Documentation** is chronically out of date and unreliable
- **Trust** in many data systems is low
- There are **many different tools** to help with this...




*Part 1:*

*The Wonderful World of (Open Source)  
Data Quality Tools in Python*





# Different aspects of data "quality"






# There are lots of different tools in the space...


- I focused on “single purpose” tools rather than end-to-end data processing packages
- It’s surprisingly hard to find a lot of open source “python data quality” packages
  - Note: The **commercial space** here is growing quickly
  - A lot of these **aren’t actively maintained**
  - I marked active projects (=some activity in past 6 months) with \*
- Let me know if I’ve missed anything!

# Running example: NYC taxi data




Number of passengers per ride in January 2019

# How do we prevent against issues like this in our data?



Number of passengers per ride in [January 2019](#)



Number of passengers per ride in [February 2019](#)

# Pure profiling tools

- **Pandas Profiling\***: Like a sophisticated extension of `.describe()` on Pandas dataframes, creates a more detailed profile report of the data.
  - <https://github.com/pandas-profiling/pandas-profiling>

----- Data profiling

----- Data documentation

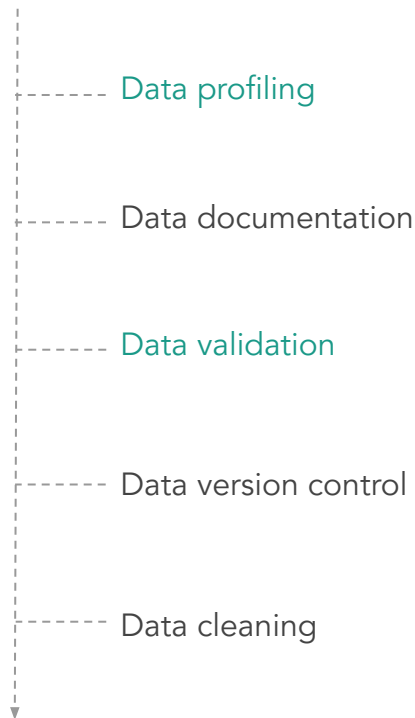
----- Data validation

----- Data version control

----- Data cleaning

# Profiling + validation

- **TDDA\***: Allows specifying and verifying constraints on data, provides some lightweight profiling
  - <https://pypi.org/project/tdda/>
- **pydqc**: Does some data profiling (show basic stats about a table), compares columns between tables, compares tables for identity.
  - <https://github.com/SauceCat/pydqc>



# Profiling, validation, and documentation

- **Great Expectations\***: Allows you to write declarative data tests ("I expect this table to have between x and y number of rows"), get validation results from those tests, and output a report that documents the current state of your data
  - <http://greatexpectations.io>

----- Data profiling

----- Data documentation

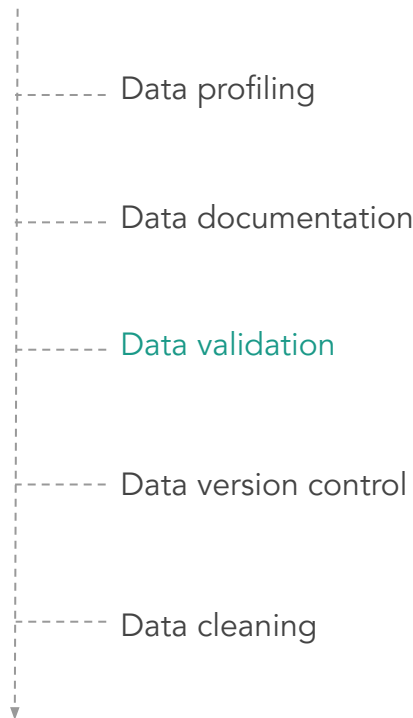
----- Data validation

----- Data version control

----- Data cleaning

# Pure data validation / testing tools

- **Bulwark\***: Data testing framework that lets you add tests on methods that return Pandas dataframes. Has some built-in tests and allows custom methods for tests.
  - <https://github.com/ZaxR/bulwark>
- **Engarde**: A precursor to Bulwark, allows you to add decorators with data assertions to methods that return Pandas data frames.
  - <https://engarde.readthedocs.io/en/latest>



# Pure data validation / testing tools

- **Voluptuous\***: Allows you to specify a "schema" to validate JSON/YAML
  - <https://github.com/alecthomas/voluptuous>
- **Opulent Pandas**: A df-focused "version" of Voluptuous
  - <https://pypi.org/project/opulent-pandas/>
- **mobydq\***: Data validation web app that allows you to check for "indicators" such as completeness, freshness, latency, validity.
  - <https://github.com/ubisoft/mobydq>

----- Data profiling

----- Data documentation

----- Data validation

----- Data version control

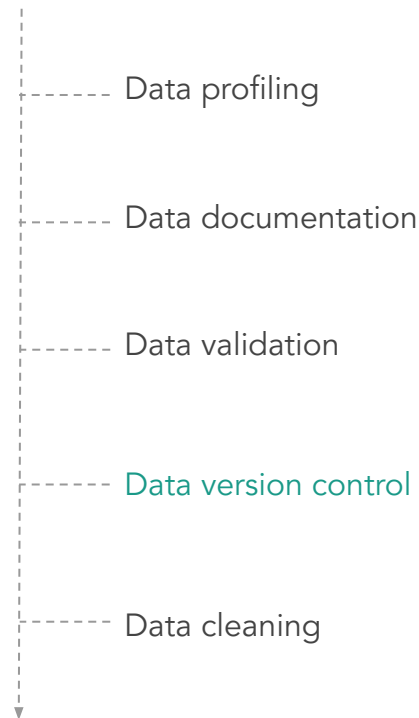
----- Data cleaning





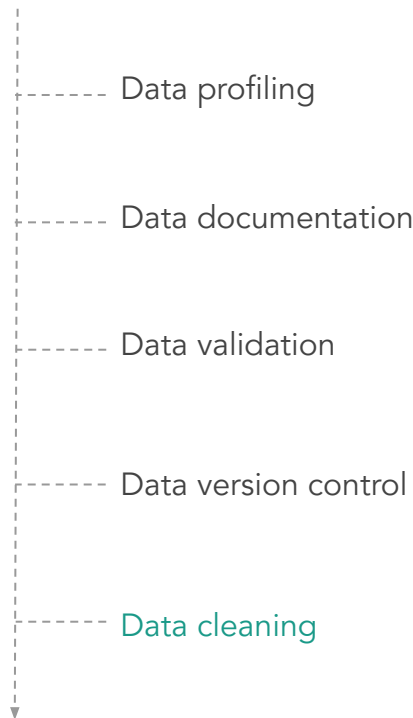
# Data version control

- **dvc\***: Built on top of git, offers version control specifically for data and Machine Learning models.
- (Any others?)



# Data cleaning

- **dedupe\***: Uses fuzzy matching to perform de-duplication and entity resolution in data.
  - <https://github.com/dedupeio/dedupe>
- **datacleaner**: Does some basic cleaning operations on data frames, e.g. dropping rows with missing values, codes strings as numeric values, etc.
  - <https://github.com/rhievery/datacleaner>
- (and many more... pyjanitor)





*Part 2:*


*Deep Dive Into Great Expectations*

*"Our stakeholders would notice data issues before we did... which really eroded trust in the data and our team"*

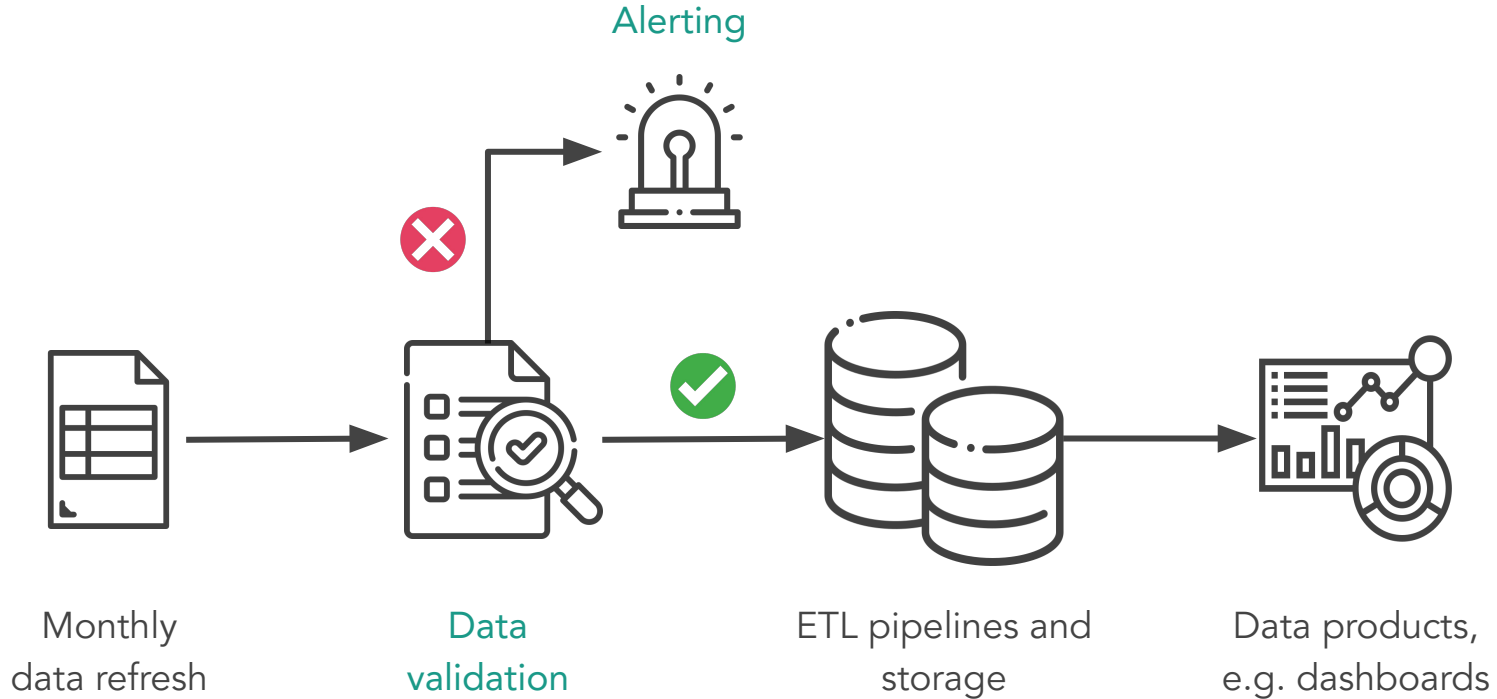


*A Great Expectations user*

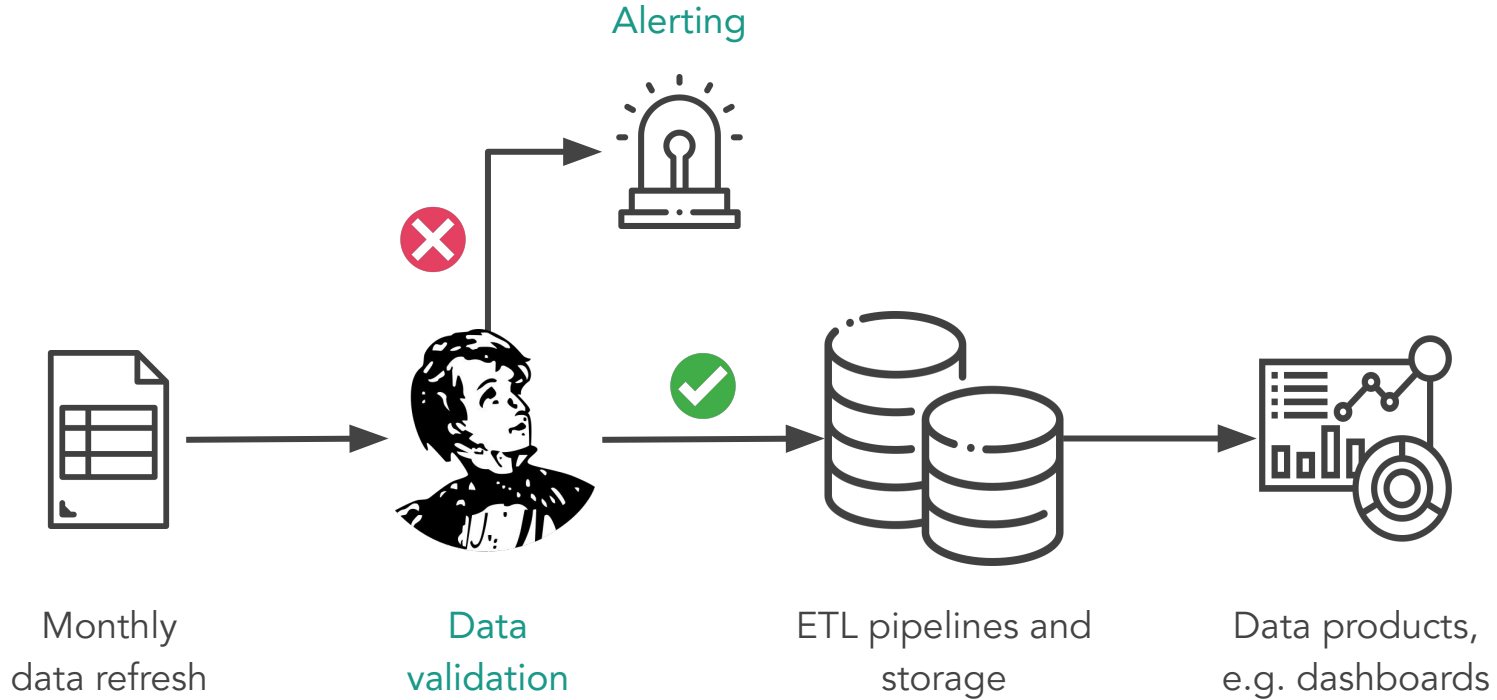
# A typical data pipeline



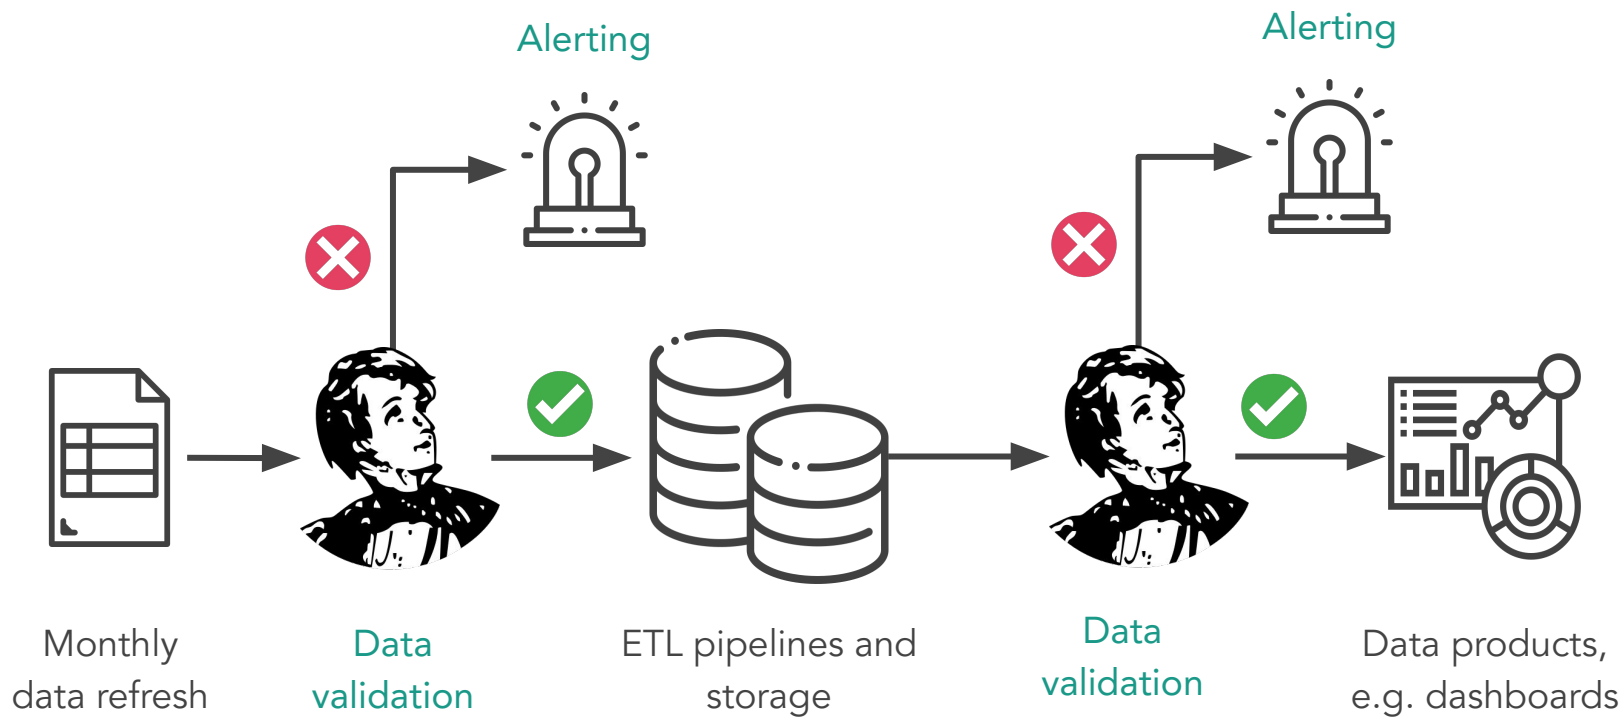
# A better data pipeline



# A better data pipeline...



# Or even better...







## Expectation Validation Result

Evaluates whether a batch of data matches expectations.

### Actions

Validation Filter:

Show All Failed Only

 How to Edit This Suite

Show Walkthrough

### Table of Contents

**Overview**

[dropoff\\_datetime](#)

[fare\\_amount](#)

[passenger\\_count](#)

### Overview

Expectation Suite: [taxi.demo](#)

Status: ✖ Failed

## We have a test suite...

### Statistics

Evaluated Expectations	8
Successful Expectations	7
Unsuccessful Expectations	1
Success Percent	87.5%

[Show more info...](#)

### Info

Great Expectations Version	0.11.9+14.g1434d72c
Run Name	20200804T163846.210132Z
Run Time	2020-08-04T16:38:46.210132+00:00

### Batch Markers

ge_load_time	20200804T163845.408935Z
--------------	-------------------------

### Batch Kwargs

data_asset_name	yellow_tripdata_staging
-----------------	-------------------------



## Actions

Validation Filter:

[Show All](#) [Failed Only](#)[✎ How to Edit This Suite](#)[Show Walkthrough](#)

## Table of Contents

[Overview](#)[dropoff\\_datetime](#)[fare\\_amount](#)[passenger\\_count](#)[pickup\\_datetime](#)[trip\\_distance](#)[vendor\\_id](#)

## passenger\_count

Status	Expectation	Observed Value
✖	<div>values must always be between 1 and 6.</div> <div>1579 unexpected values found. ≈15.79% of 10000 total rows.</div> <div>Sampled Unexpected Values</div> <div>0.0</div>	≈15.79% unexpected

... in which we express what we "expect" from our data

## pickup\_datetime

Status	Expectation	Observed Value
✔	values must never be null.	100% not null

## trip\_distance

# What is Great Expectations?

```
> pip install great_expectations
```

It's an open source Python library

```
my_data_project
├── great_expectations
│   ├── checkpoints
│   ├── expectations
│   │   └── taxi
│   │       └── demo.json
│   ├── plugins
│   ├── uncommitted
│   └── great_expectations.yml
```

That operates by creating tests in code and storing a collection of them as a "suite" in JSON

# What is an Expectation?

```
expect_column_values_to_be_between(  
    column='passenger_count',  
    min_value=1,  
    max_value=6  
)
```

It's a statement about what we expect from our data, expressed as a method in Python

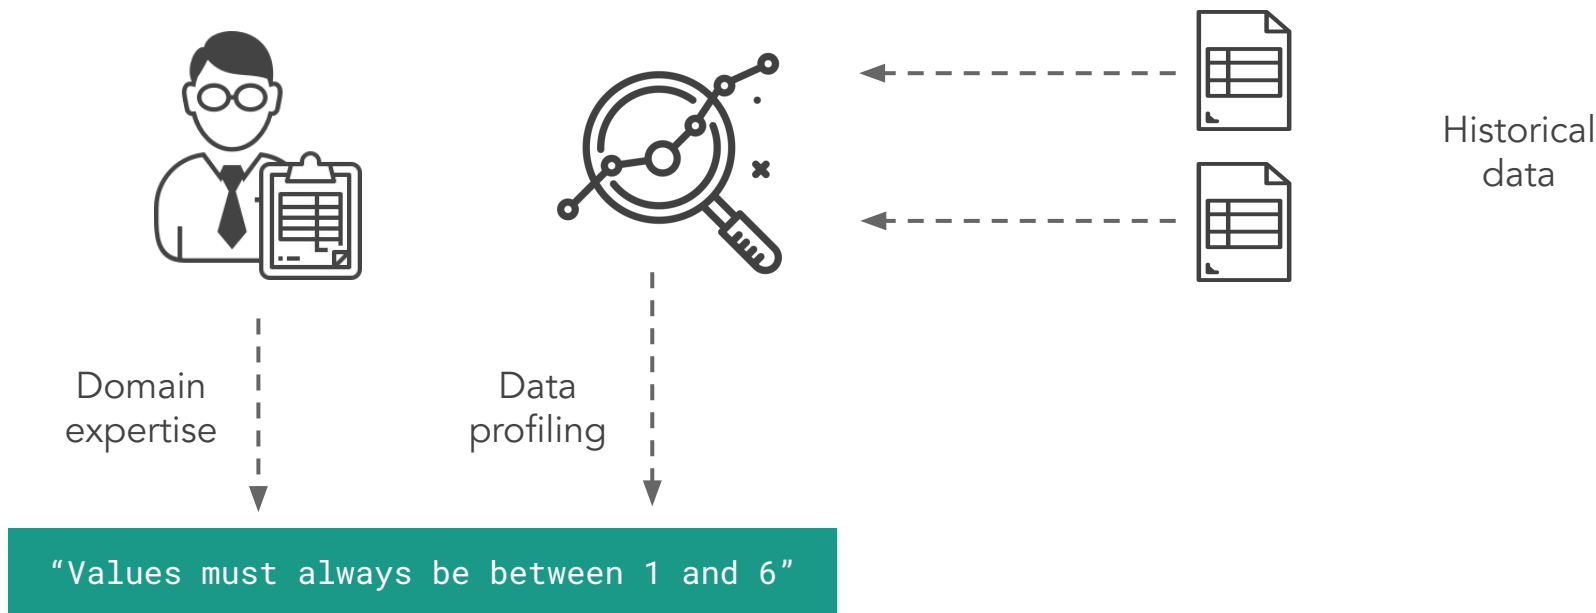
```
{  
  "expectation_type": "expect_column_values_to_be_between",  
  "kwargs": {  
    "column": "passenger_count",  
    "min_value": 1,  
    "max_value": 6  
  },  
}
```

That is stored in JSON

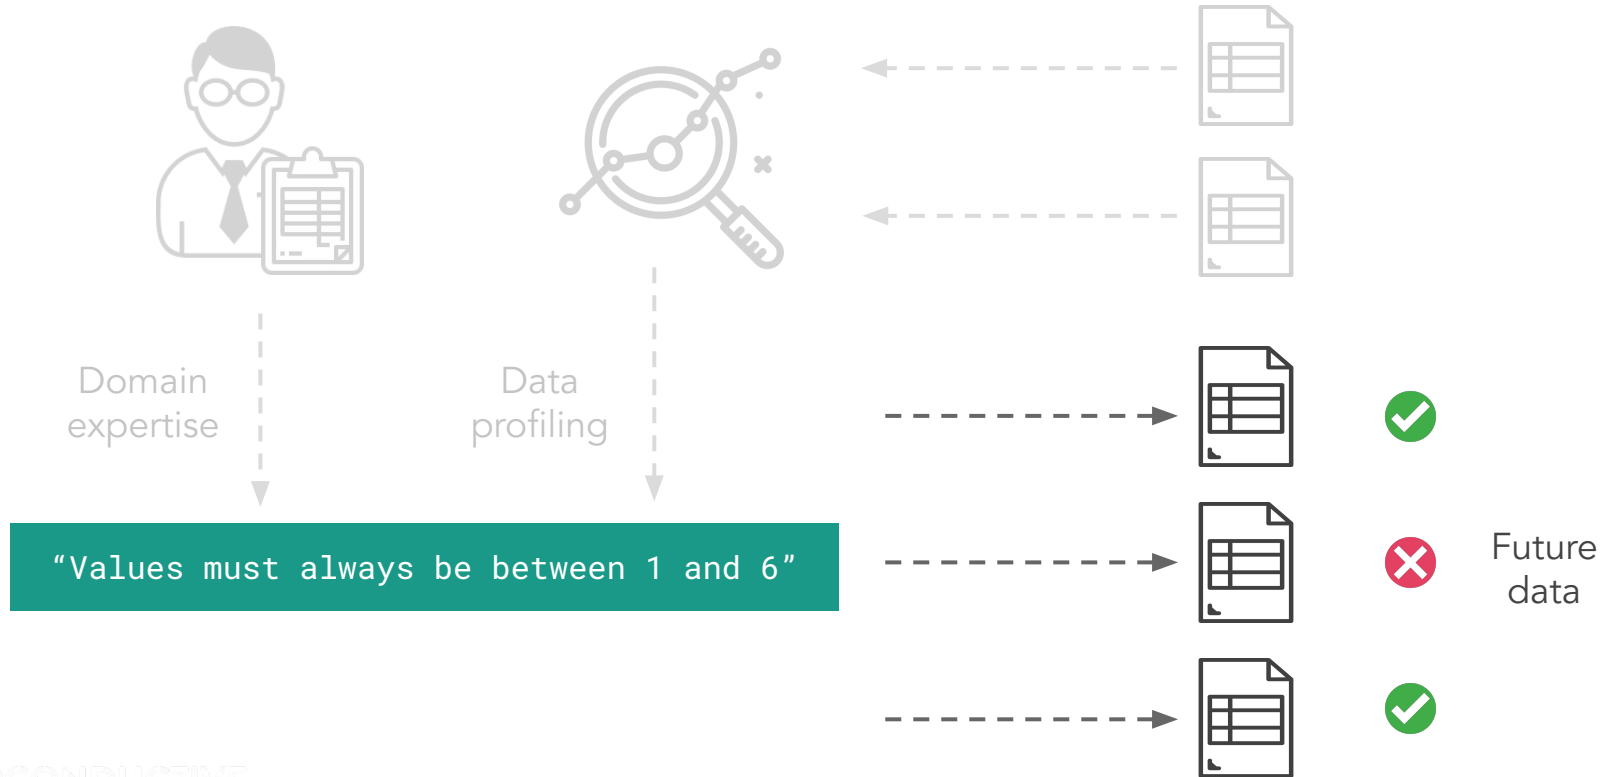
"Values must always be between 1 and 6"

And can be translated into a human-readable format

# Built-in profiling to “scaffold” Expectations



# Validation of new data using Expectations



# Different types of Great Expectations workflows

## Interactive workflows

Import the GE library in a notebook (local, Databricks, etc.)

Connect to a batch of data

Profile and validate data on-the-fly and iteratively using Expectations

## Batch processing workflows

Create and store Expectation Suites

Load Expectation Suites to validate incoming batches of data

Validation results determine alerting and pipeline behavior

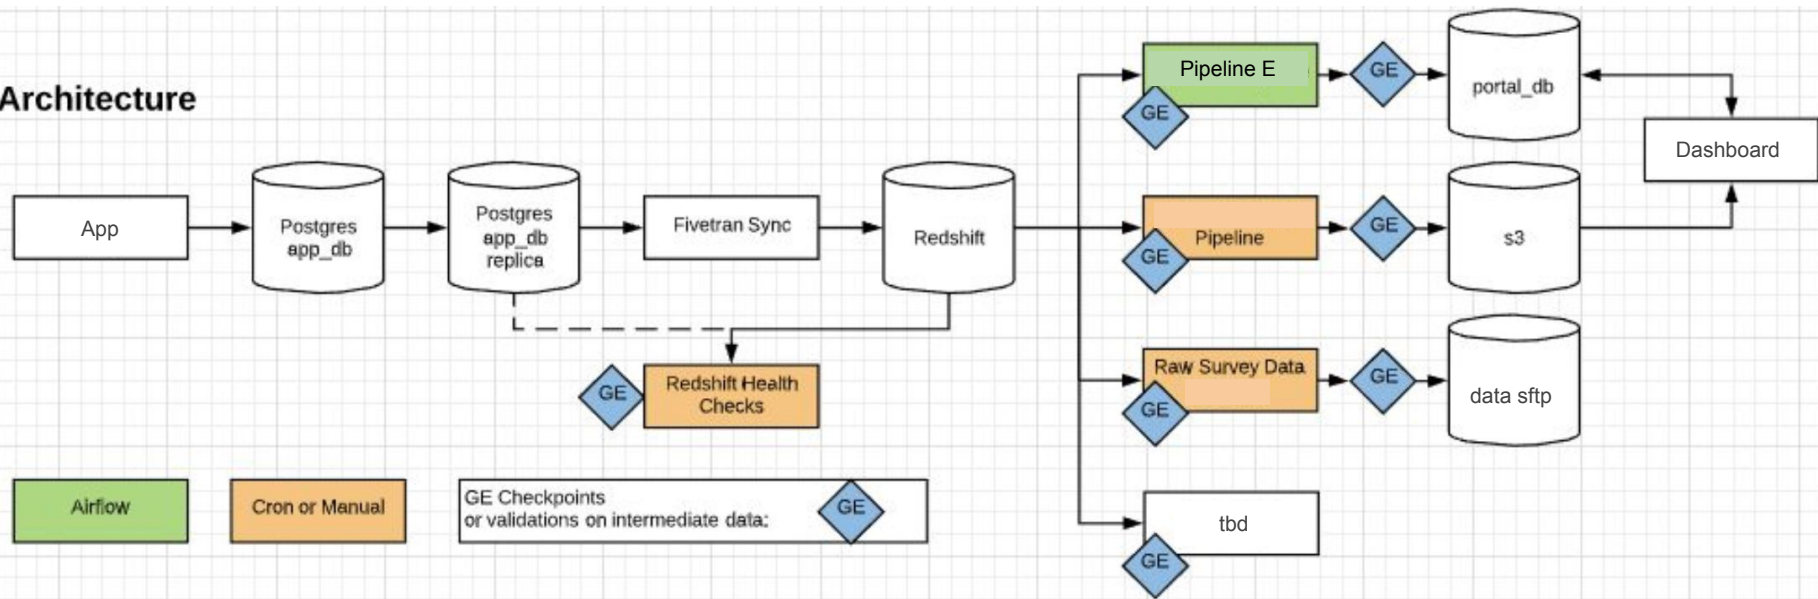
## CI/CD workflows

Local tests of pipeline changes against input fixtures during code development process

Automated testing against input fixtures, e.g. when opening a PR

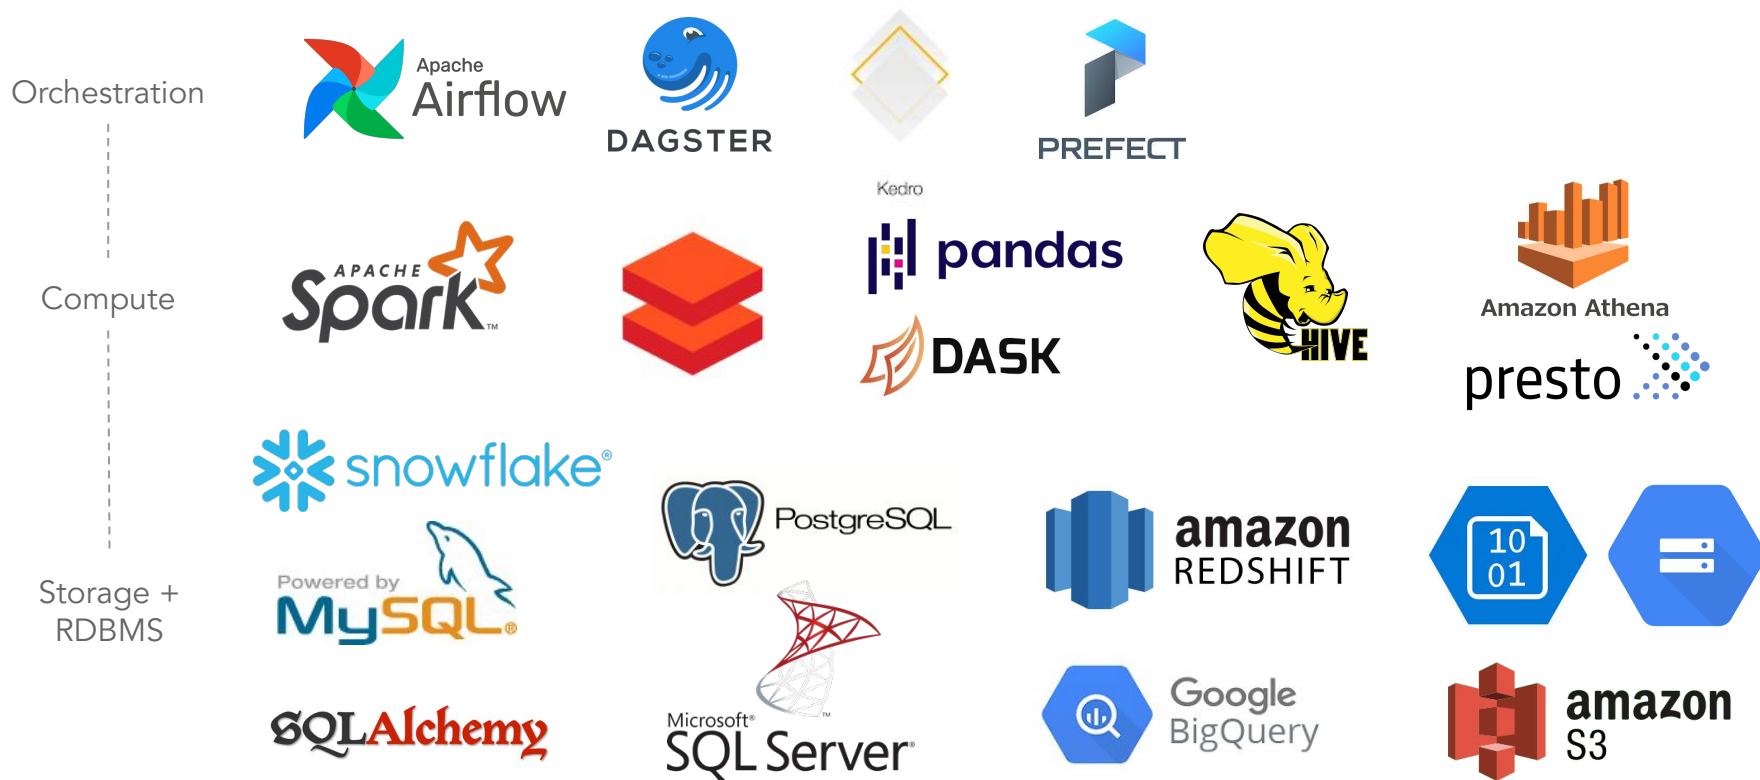
# Sample architecture

## Architecture





# Integrations





*Part 3:*

*Demo time! "Let's do it live"*

# In this demo, I'll show you a quick workflow

- Initialize a data context
- Connect to a Datasource
- Create an Expectation Suite with an automated profiler
- (If time allows: Create a Checkpoint to validate new data)
- Navigate Data Docs



*That's it! Questions?*  
*Sam Bail @spbail*





## Great Expectations

Join our Slack channel!  
[greatexpectations.io/slack](https://greatexpectations.io/slack)